

ECE 253

Digital Image Processing

Introduction to Convolutional Neural Networks

HW 3/4 Problem:

Given an image, can we determine the object?





Object: Stop Sign

HW 3/4 Problem:

Given an image, can we determine the object?

Object Detection by Correlation



template

How do we locate the "template" (Object) in the image?

- Minimize
$$E(i, j) = \sum_m \sum_n [f(m, n) - t(m - i, n - j)]^2$$
$$= \sum_m \sum_n [f^2(m, n) + t^2(m - i, n - j) - 2f(m, n)t(m - i, n - j)]$$
- Maximize
$$R_{ij}(i, j) = \sum_m \sum_n t(m - i, n - j)f(m, n)$$

Cross-correlation



Object: Stop Sign

HW 3/4 Problem:

Given an image, can we determine the object?

Why can't we just use correlation?



Object: Stop Sign

HW 3/4 Problem:

Given an image, can we determine the object?

Why can't we just use correlation?



Object: Stop Sign

HW 3/4 Problem:

Given an image, can we determine the object?

Why can't we just use correlation?



Object: Stop Sign



Many configurations (orientation, lighting, context) = too much information to be accurately and robustly contained by a single template filter, or even a group of filters.

Plus, many other possible objects!

Neural networks can be used to solve problems like **classification**, **clustering**, **regression**, and **generation** by **learning weights** as a way of **approximating a function** F which maps input data to the desired output.


Input
Data



$F(x)$




Predicted
Labels
(for classification
problems)






Modeled in Biology

At the junction between two neurons (synapse), an action potential causes a neuron to release a chemical neurotransmitter (after voltage passes action potential threshold).

Neuron







Neural Networks in Code

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```


Neural Networks in Code

```
class Net(nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(3, 6, 5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.conv2 = nn.Conv2d(6, 16, 5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)
```


```
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = x.view(-1, 16 * 5 * 5)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x
```



```
net = Net()
```

What is a linear layer?





“Linear”, “Fully-Connected”, “Dense” Layer








Where is the action potential response??



Perceptron





```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x



net = Net()
```

What is a linear layer?

What is a “relu”? This is a type of **Activation Layer**.

Activation Functions


- Activation function allows for non-linearity in the network approximation of function F .
- This is important since most complex data we collect is not linear!



- Pitfall: by reducing input to binary output, we lose information. (Recall Quantization exercises)

Activation Functions

Continuous output (instead of binary) helps provide more information from each neuron.




$$1 + \exp(-\sum_j w_j x_j - b)$$

Activation Functions

Continuous output (instead of binary) helps provide more information from each neuron.


Sigmoid



1

$$1 + \exp(-\sum_j w_j x_j - b)$$

ReLU



```

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()

```

What is a **convolutional layer**?

What is a linear layer?

What is a “relu”? This is a type of **Activation Layer**.

Modeling for 2D Spatial Input (i.e. images)

- Assume there is a spatial relationship between data points (the neighborhood of a point contains useful information).
- How can we extract that information in an efficient way?

Terminology

Convolutional Layer


1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved
Feature

Convolutional Layer



Convolutional Layer

32x32x3 image



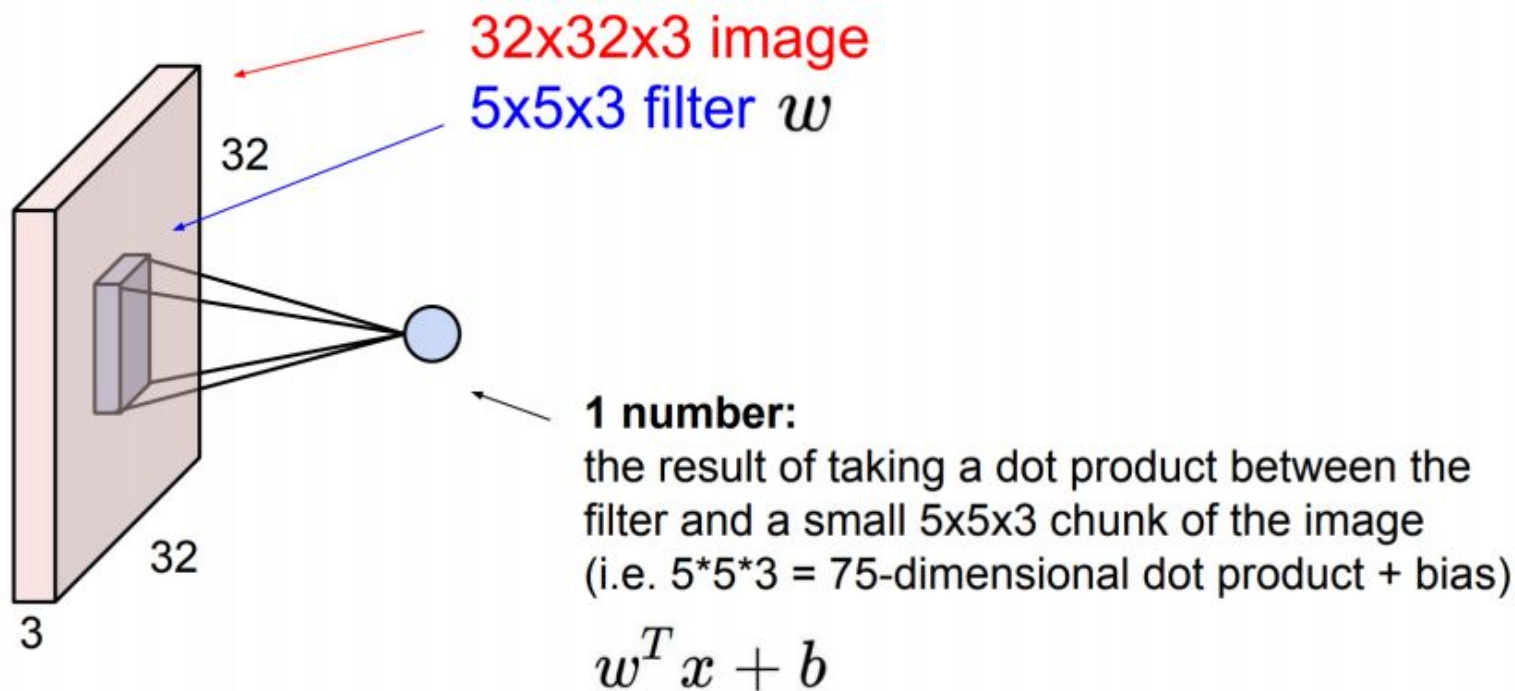
5x5x3 filter




Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

* Ok, technically this is correlation, but “Correlational Neural Network” doesn’t sound as cool.

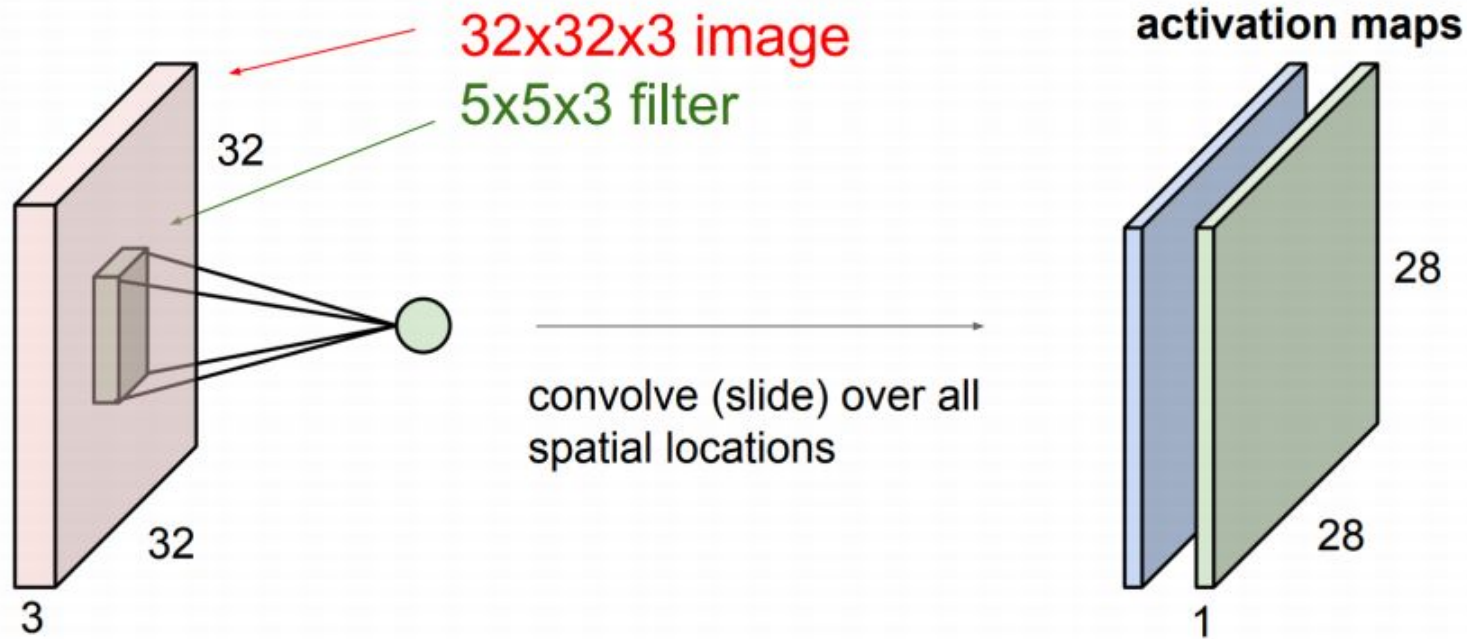
Convolutional Layer



Convolutional Layer




Convolutional Layer



Convolutional Layer

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!


```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```


What is a **convolutional layer**?

What is max pool? This is a type of **Pooling Layer**.


What is a linear layer?

What is a “relu”? This is a type of **Activation Layer**.

Max Pooling




Max Pooling




Max Pooling

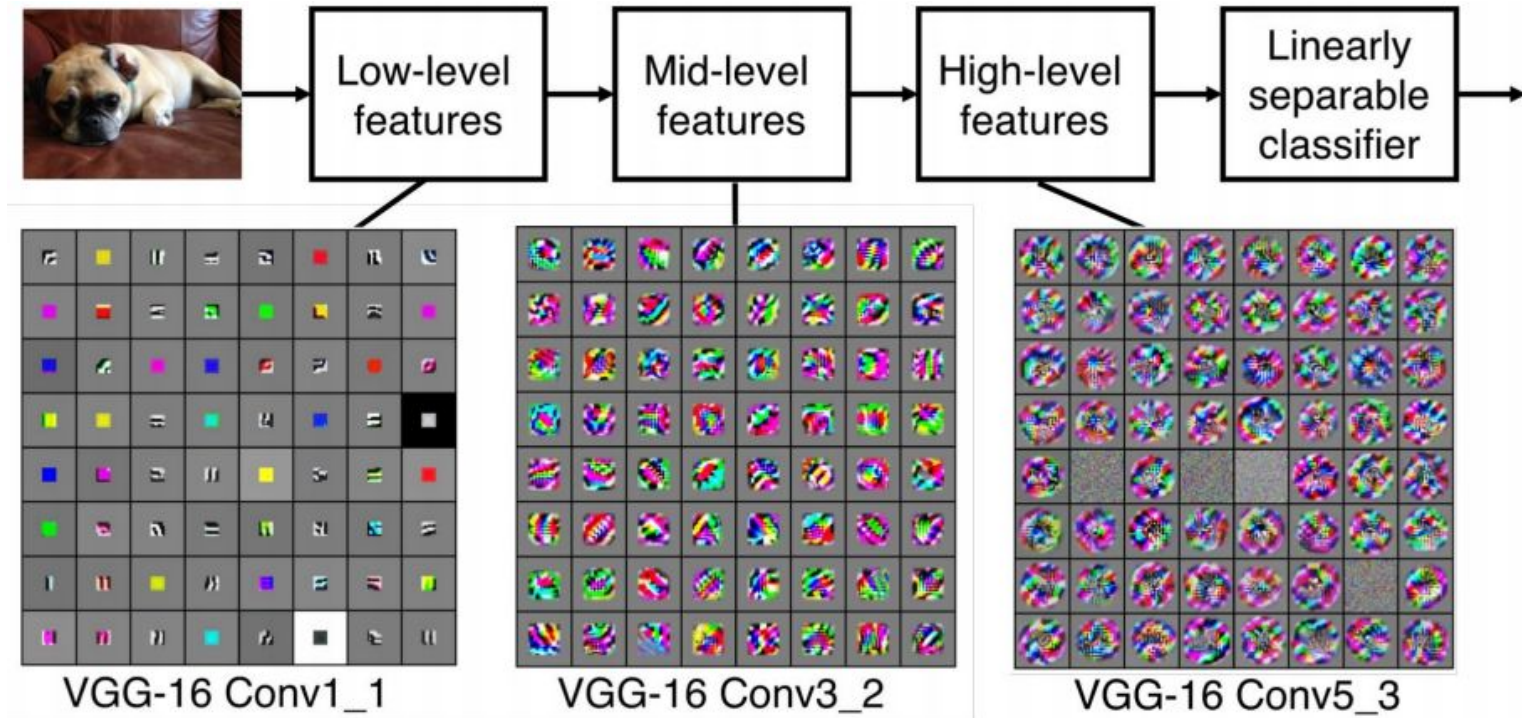
- makes the representations smaller and more manageable
- operates over each activation map independently:



Putting it all together



Putting it all together



```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
net = Net()
```

What is a **convolutional layer**?

What is max pool? This is a type of **Pooling Layer**.

What is a linear layer?

What is a “relu”? This is a type of **Activation Layer**.

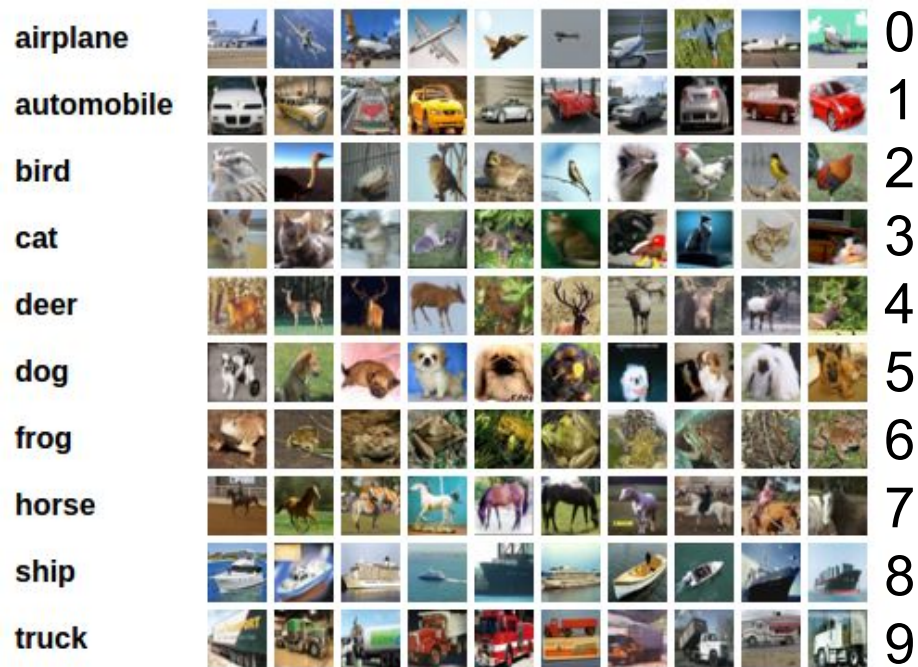
What does the output (x) look like?

We wanted something like “Stop Sign”...

How do we turn a name into a vector?

“One-Hot” Encoding
































































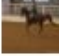


























First, predefine:



How do we turn a name into a vector?

“One-Hot” Encoding

First, predefine:

airplane										0
automobile										1
bird										2
cat										3
deer										4
dog										5
frog										6
horse										7
ship										8
truck										9



$$\text{Dog} \rightarrow \text{Class 5} \rightarrow \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$





























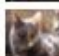



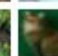


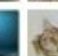







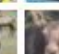



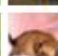




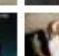










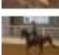
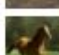


















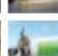



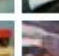




$$\text{Cat} \rightarrow \text{Class 3} \rightarrow \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

How do we turn a name into a vector?

“One-Hot” Encoding

First, predefine:

airplane										0
automobile										1
bird										2
cat										3
deer										4
dog										5
frog										6
horse										7
ship										8
truck										9



Dog → Class 5


$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$




Cat → Class 3

$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

Forward Propagation



Forward Propagation



- If we initialize all weights (w) randomly, there will be a difference between the predicted output (y_{pred}) and true, expected class (y).
- We quantify this difference via a **loss function** (L).

Loss Functions

Mean Squared-Error

$$\frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

- * n is the number of data points
- * Y_i represents observed values
- * \hat{Y}_i represents predicted values

Loss Functions

Mean Squared-Error

$$\frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

- * n is the number of data points
- * Y_i represents observed values
- * \hat{Y}_i represents predicted values

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions



Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions



Dog: 0.00

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions



Dog: 0.00
Cat: 0.02

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions



Dog: 0.00
Cat: 0.02
Fish: 0.00

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions




Dog: 0.00
Cat: 0.02
Fish: 0.00
Guacamole: 0.98

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- \log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions



Dog: 0.00
Cat: 0.02
Fish: 0.00
Guacamole: 0.98


y	p	-ylog(p)
0	0	0
1	0.02	Large, positive
0	0	0
0	0.98	0

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

- M - number of classes (dog, cat, fish)
- log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c

Loss Functions




Dog: 0.00
Cat: 0.02
Fish: 0.00
Guacamole: 0.98

y	p	-ylog(p)
0	0.02	0
1	0.98	Near 0
0	0	0
0	0	0

Cross-Entropy

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$


- M - number of classes (dog, cat, fish)
- log - the natural log
- y - binary indicator (0 or 1) if class label c is the correct classification for observation o
- p - predicted probability observation o is of class c



How do we choose these weights?

Back Propagation (Optimization)

- Ideally, $L = 0$ (predicted output matches correct output)




Idea: reduce loss by adjusting the weights!

Back Propagation (Optimization)

- Ideally, $L = 0$ (predicted output matches correct output)
- Is this always possible?

Training objective:

$$W^* = \operatorname{argmin}_W \sum_i L(f_W(X^{(i)}), Y^{(i)})$$



Back Propagation (Optimization)

- We can reduce the loss by changing the values of the weights.
- How severely should we change the weight?

*Note, these equations refer to 'L' (loss) as 'C' (cost).

Learning Rate



$$W^+ = W - \eta \nabla C$$


$$\nabla C = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \vdots \\ \frac{\partial C}{\partial w_n} \end{bmatrix}$$

Back Propagation (Optimization)

Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$




Back Propagation (Optimization)

Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = \quad , \quad \frac{\partial q}{\partial y} =$$




Back Propagation (Optimization)

Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$



Back Propagation (Optimization)


Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = , \frac{\partial f}{\partial z} =$$



Back Propagation (Optimization)


Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$



Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial y}$$

Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

$$\frac{\partial f}{\partial y}$$

Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial x}$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Back Propagation (Optimization)

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial x}$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Back Propagation (Optimization)

- We can reduce the loss by changing the values of the weights.
- How severely should we change the weight?

*Note, these equations refer to 'L' (loss) as 'C' (cost).

Learning Rate



$$W^+ = W - \eta \nabla C$$

$$\nabla C = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \vdots \\ \frac{\partial C}{\partial w_n} \end{bmatrix}$$

Back Propagation (Optimization)

From Professor Trivedi:

UCSD interesting Trivia:

Back propagation was “invented” in 1986...

David E. Rumelhart, Geoffrey E. Hinton und Ronald J. Williams.
Learning representations by back-propagating errors., Nature
(London) 323, S. 533-536.

All three authors were working at UCSD in 1986!

Back Propagation (Optimization)

- We can reduce the loss by changing the values of the weights.
- How severely should we change the weight?

*Note, these equations refer to 'L' (loss) as 'C' (cost).

$$W^+ = W - \eta \nabla C$$

Learning Rate



$$\nabla C = \begin{bmatrix} \frac{\partial C}{\partial w_1} \\ \frac{\partial C}{\partial w_2} \\ \vdots \\ \frac{\partial C}{\partial w_n} \end{bmatrix}$$

Can be extremely slow and memory inefficient ...


- Use smarter update rules that keep memory of past gradient values (eg. Momentum, adagrad, Adam)
- Use minibatches of data to update gradients, rather than the entire dataset

Training a Neural Network

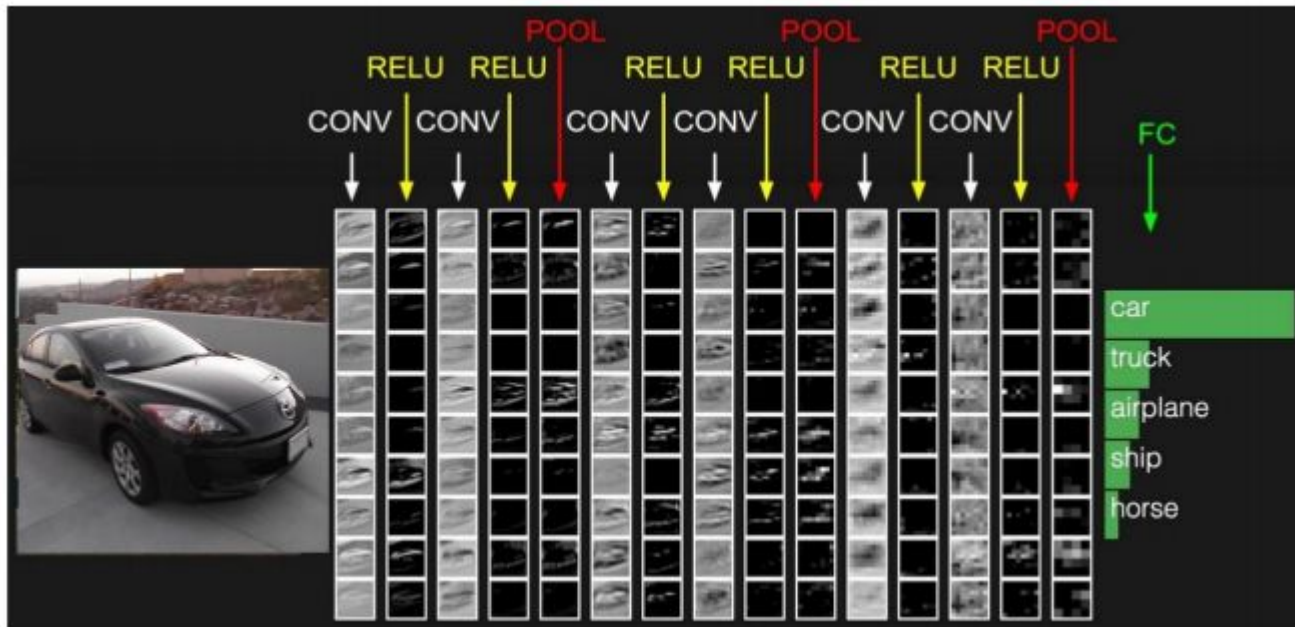
- Choose a network structure
 - Shape of input/output?
 - How many hidden layers, and how should the layers be connected?
- Initialize the network weights
- For training data (x, y) :
 - Forward pass data x through the network to make a prediction, y_{pred}
 - Calculate loss $L(y, y_{\text{pred}})$
 - Change weights according to learning rate & gradient of loss with respect to weights
- Repeat ... and repeat ... and repeat
- For new data points (x) without a known output (y) , we can pass x through the trained network to form a prediction!

When should we stop training?

- Each time we loop through the dataset is called an 'epoch'.
- Divide data into a *training set* and *validation set*.
 - Do not train on validation set.
 - Stop training when the validation error begins to increase.
 - Training beyond this is referred to as 'overfitting,' since the model becomes overfit to the training data (which is only a subset of the possible real data!).



Quick Recap:



HW 4 Problem:

Given an image, can we determine an object label for all pixels?
("Segmentation")



Segmentation Challenges

Based on what we have seen for image classification, how would you go about creating a network for pixel classification (i.e. image segmentation)?

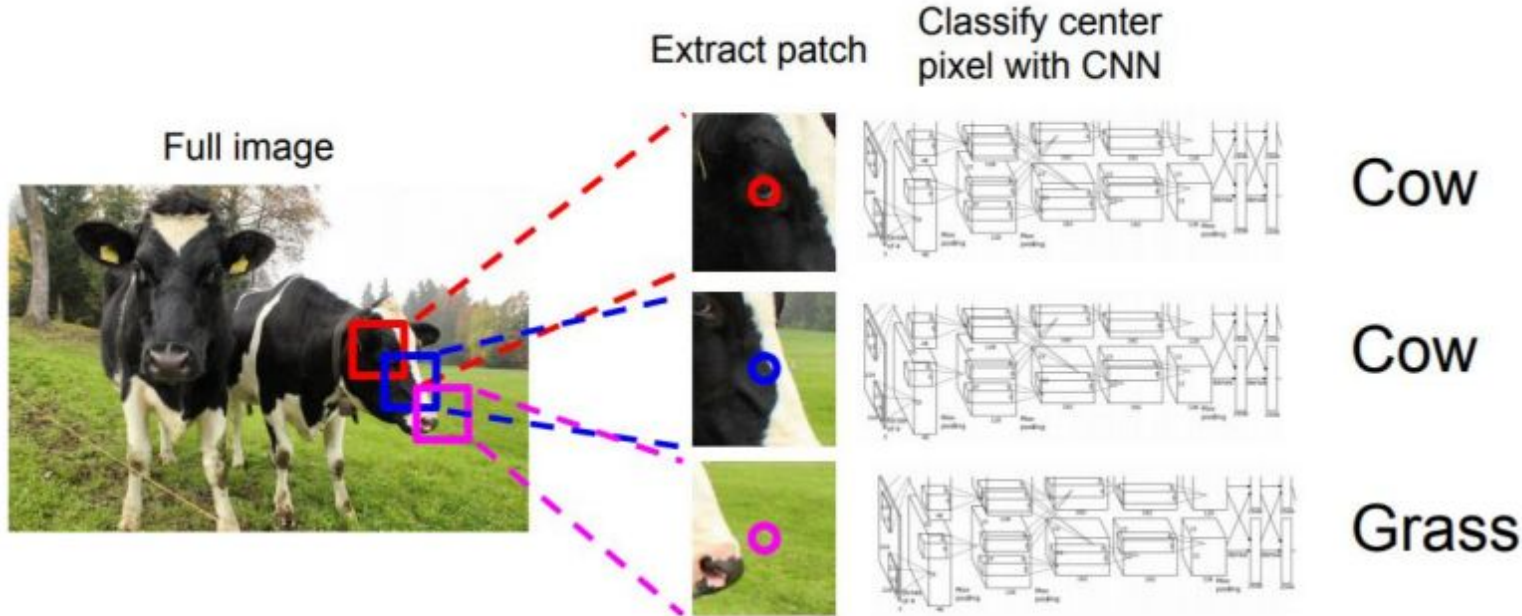
Segmentation Challenges

Based on what we have seen for image classification, how would you go about creating a network for pixel classification (i.e. image segmentation)?

Preliminary idea (Ciresan et al.):

Move a sliding window centered over each pixel. For each window location, predict the class of the pixel.

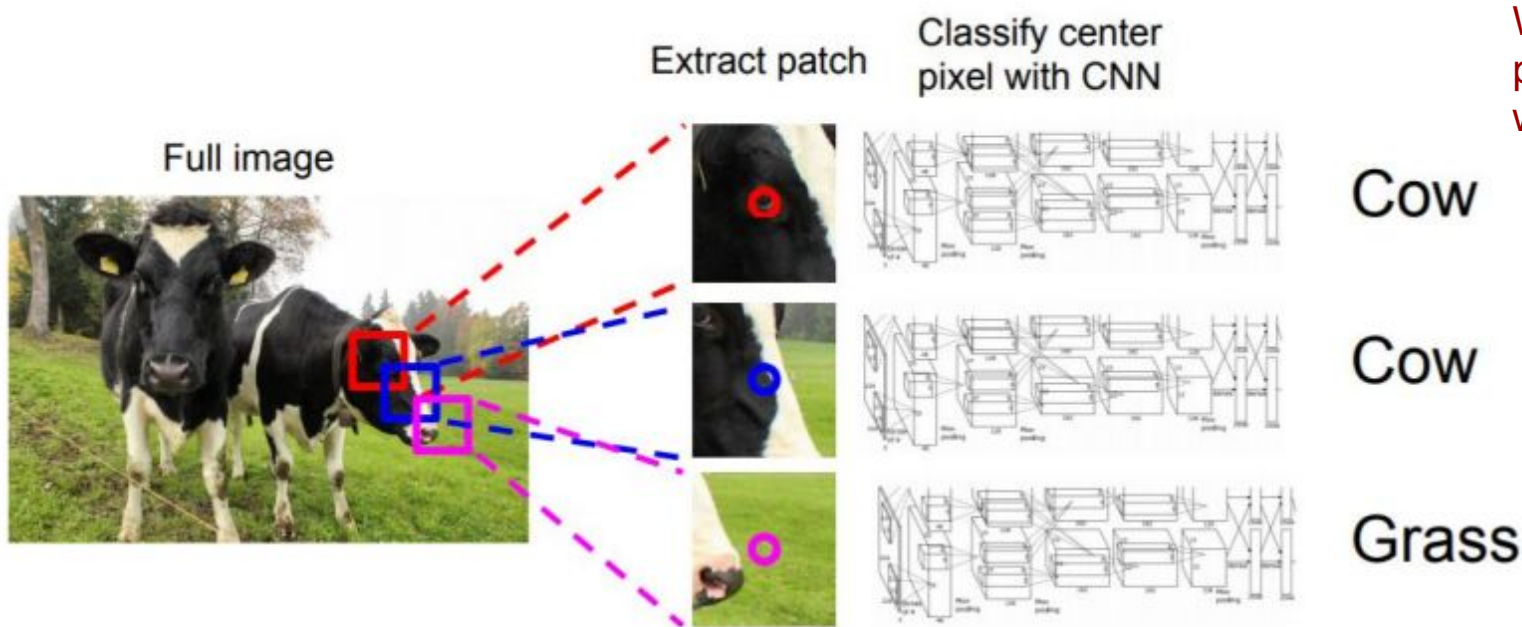
Semantic Segmentation Idea: Sliding Window



Farabet et al, "Learning Hierarchical Features for Scene Labeling," TPAMI 2013

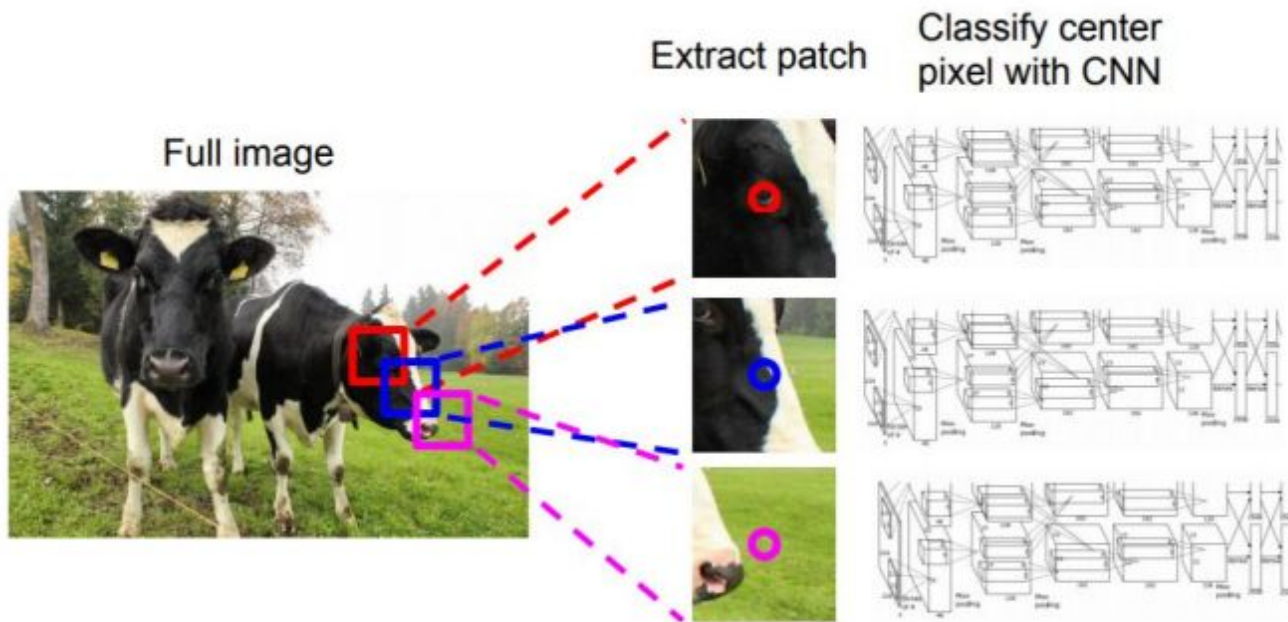
Pinheiro and Collobert, "Recurrent Convolutional Neural Networks for Scene Labeling", ICML 2014

Semantic Segmentation Idea: Sliding Window



What are some possible problems with this method?

Semantic Segmentation Idea: Sliding Window



What are some possible problems with this method?

Cow

- (1) Inefficient
- (2) Distinguishing features occur at different scales


Cow

Grass

Addressing Inefficiency

Semantic Segmentation Idea: Fully Convolutional

Design a network as a bunch of convolutional layers to make predictions for pixels all at once!




Addressing Feature Scales

Semantic Segmentation Idea: Fully Convolutional

Design network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!



Input:
 $3 \times H \times W$



Predictions:
 $H \times W$

Upsampling

Nearest Neighbor

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

“Bed of Nails”

1	2
3	4

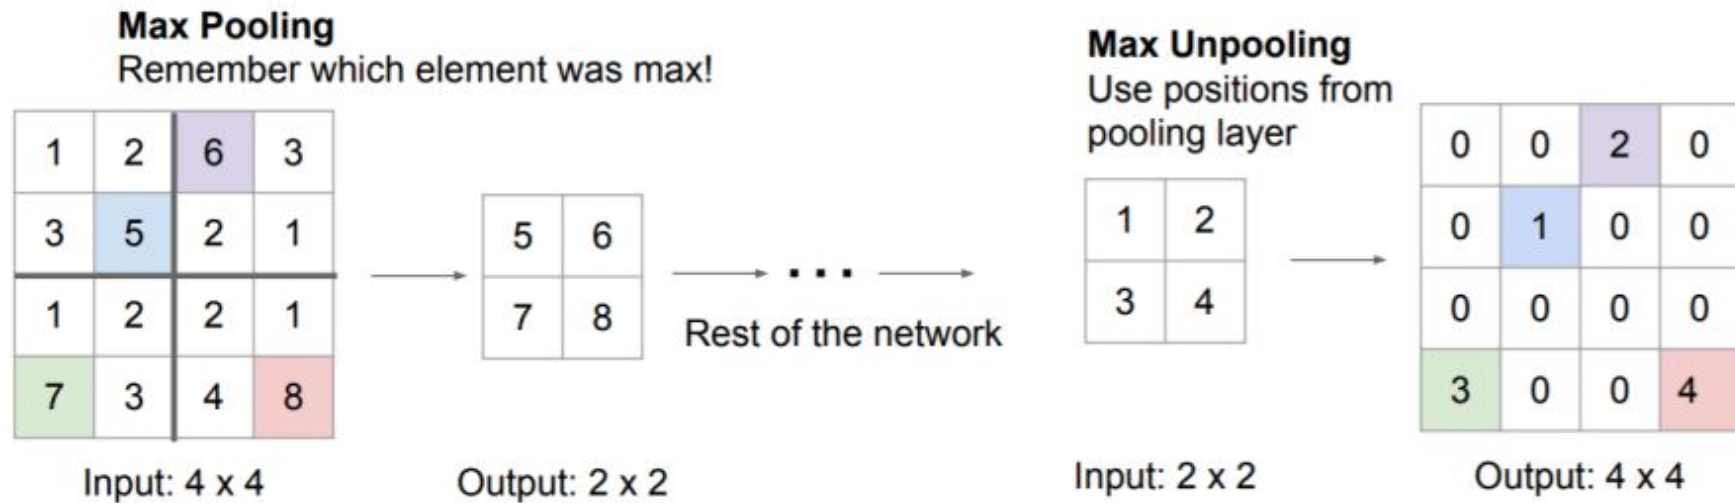


1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0


Input: 2 x 2

Output: 4 x 4


Upsampling




Corresponding pairs of
downsampling and
upsampling layers




Fully Convolutional Networks



Fully Convolutional Networks



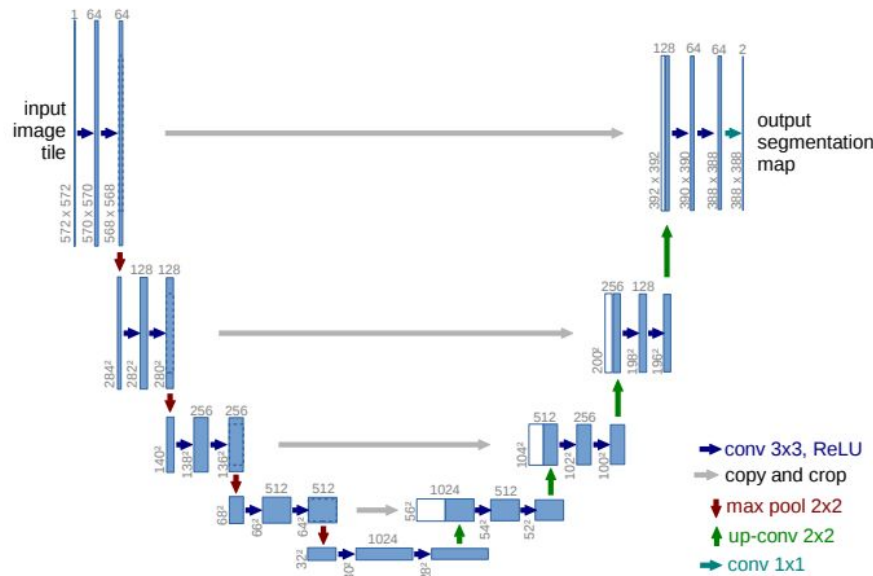
Fully Convolutional Networks



Output size matches input!

Refining the Approach: U-Net

“The architecture consists of a contracting path to capture context and a symmetric expanding path that enables precise localization.” (Ronneberger et al.)



Refining the Approach: U-Net

